

# CS 583: PROBABILISTIC GRAPHICAL MODELS

## TOPIC: STRUCTURE LEARNING



**Mustafa Bilgic**



<http://www.cs.iit.edu/~mbilgic>



<https://twitter.com/bilgicm>

# TASK

- We are given fully observed data
- Our task is to learn the structure (and the parameters)

# GOAL: KNOWLEDGE DISCOVERY

- There are at least two problems
  - Unidentifiability
    - Given observational data, the best we can hope to learn is an I-Equivalent structure to  $G^*$
  - Weak connections
    - Due to noise, two variables will almost never be independent
    - If a connection is weak, how do we determine if it is due to noise or it is a real weak connection?

# GOAL: DENSITY ESTIMATION

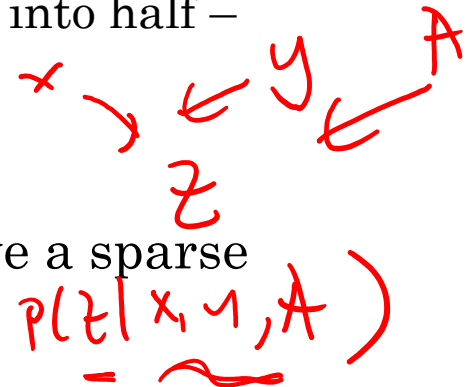
- Adding an unnecessary edge is very costly for parameter estimation

- Each unnecessary binary parent divides the data into half – data fragmentation

- The estimated parameters have high variance

- For better generalization, it is preferable to have a sparse structure

- In fact, if the data is limited, it might be even better to have a sparser structure than  $G^*$*



# BAYESIAN NETWORKS

# THREE APPROACHES

## 1. *Constraint-based structure learning*

- View Bayesian network as a representation of independencies
- Test for conditional dependence and independence in the data
- Sensitive to failure in individual independence tests

## 2. *Score-based structure learning*

- Hypothesis space of possible models
- Find the highest scoring structure
- The hypothesis space is super exponential  $2^{(O(n^2))}$

## 3. *Bayesian model averaging*

- Rather than learning a structure, learn an ensemble of structures
- Again, the number of structures is immense
- It is efficient for some while it has to be approximate for others

# 1. CONSTRAINT-BASED APPROACHES

- In the BN slide deck, we have seen two procedures
  1. Find minimal I-MAP
  2. Find all I-equivalent structures
- We assumed that, given a distribution  $P$ , we could answer any independence question perfectly
- We'll now assume that we don't have access to  $P$  but instead we have access to a dataset  $D$

# INDEPENDENCE IN THE DATA?

- Assume we toss two coins independently
- We first toss one coin, record it, and then toss the next coin, and record the outcome
- We repeat this process 100 times
- Assume we get 25 H/H, 23 H/T, 27 T/H, 25 T/T
- Based on the data, can we say that the two coins are independent?



# INDEPENDENCE TESTS

$$\hat{P}(x, y) = \hat{P}(x) \hat{P}(y)$$

- $\chi^2$  statistic

$$d_{\chi^2}(D) = \sum_{x,y} \frac{\left( M[x, y] - M \times \hat{P}(x) \times \hat{P}(y) \right)^2}{M \times \hat{P}(x) \times \hat{P}(y)}$$

$$M[x, y] = M \cdot \hat{P}(x, y)$$

- Mutual information

$$d_I(D) = \sum_{x,y} \hat{P}(x, y) \log \frac{\hat{P}(x, y)}{\hat{P}(x) \hat{P}(y)}$$

$$\hat{P}(x, y) = \hat{P}(x) \hat{P}(y)$$

- $\chi^2$  and MI are zero when X and Y are independent
- Extension of  $\chi^2$  statistic to conditional independence queries

$$d_{\chi^2}(D) = \sum_{x,y,z} \frac{\left( M[x, y, z] - M \times \hat{P}(z) \times \hat{P}(x | z) \times \hat{P}(y | z) \right)^2}{M \times \hat{P}(z) \times \hat{P}(x | z) \times \hat{P}(y | z)}$$



# THREE APPROACHES

## 1. *Constraint-based structure learning*

- View Bayesian network as a representation of independencies
- Test for conditional dependence and independence in the data
- Sensitive to failure in individual independence tests

## 2. *Score-based structure learning*

- Hypothesis space of possible models
- Find the highest scoring structure
- The hypothesis space is super exponential  $2^{(O(n^2))}$

## 3. *Bayesian model averaging*

- Rather than learning a structure, learn an ensemble of structures
- Again, the number of structures is immense
- It is efficient for some while it has to be approximate for others

## 2. SCORE-BASED APPROACHES

- 2.1 Maximum likelihood score
- 2.2. Bayesian score

## 2.1 MAXIMUM LIKELIHOOD SCORE

- Find a model,  $\langle \mathcal{G}, \theta_{\mathcal{G}} \rangle$  that would make the data  $D$  as probable as possible

$$\begin{aligned}\max_{\mathcal{G}, \theta_{\mathcal{G}}} L(\langle \mathcal{G}, \theta_{\mathcal{G}} \rangle : D) &= \max_{\mathcal{G}} \left[ \max_{\theta_{\mathcal{G}}} L(\langle \mathcal{G}, \theta_{\mathcal{G}} \rangle : D) \right] \\ &= \max_{\mathcal{G}} L(\langle \mathcal{G}, \hat{\theta}_{\mathcal{G}} \rangle : D)\end{aligned}$$

- In order to find the maximum likelihood pair  $\langle \mathcal{G}, \theta_{\mathcal{G}} \rangle$ 
  - We should find the structure  $\mathcal{G}$  that achieves the highest likelihood when we use the MLE parameters for  $\mathcal{G}$

$$\text{score}_L(\mathcal{G} : D) = l(\hat{\theta}_{\mathcal{G}} : D)$$

## TWO VARIABLE CASE

- Given two variables  $X$  and  $Y$
- Consider two structures:  $G_{\emptyset}$  and  $G_{X \rightarrow Y}$ .
- What is  $\text{score}_L(G_{X \rightarrow Y} : D) - \text{score}_L(G_{\emptyset} : D)$ ?  $= M \cdot I(X; Y)$
- The general equation for  $\text{score}_L(G : D)$  is as follows. Proof is left as an exercise.

$$\text{score}_L(G : D) = M \sum_{i=1}^n I_{\hat{P}}(X_i; Pa_{X_i}^G) - M \sum_{i=1}^n H_{\hat{P}}(X_i)$$

# LIMITATIONS OF ML SCORE

- Can  $\text{score}_L(G_{X \rightarrow Y} : D)$  be smaller than  $\text{score}_L(G_{\emptyset} : D)$  ever?
- They are equal only when  $X$  and  $Y$  are truly independent and otherwise  $\text{score}_L(G_{X \rightarrow Y} : D)$  is always bigger
- Exercise: Show that  $I_P(X; Y \cup Z) \geq I_P(X; Y)$
- Adding a parent can never decrease the maximum likelihood score
- The highest scoring structure is the complete graph
- One remedy is to bound the indegree of the graph

## 2.2 BAYESIAN SCORE

- Bayesian approach puts a distribution on anything we are uncertain about
  - In parameter estimation, we put a distribution on  $\theta$
- In this case, we have the following distributions
  - $P(\mathcal{G})$
  - $p(\theta_{\mathcal{G}} | \mathcal{G})$
- Given data  $D$ , the posterior  $P(\mathcal{G} | D)$  is
  - $P(\mathcal{G} | D) = P(D | \mathcal{G}) P(\mathcal{G}) / P(D)$
- $P(D)$  is a normalizing constant
- $\text{score}_B(\mathcal{G} : D) = \log P(D | \mathcal{G}) + \log P(\mathcal{G})$

# $P(D | \mathcal{G})$

$$P(D | \mathcal{G}) = \int_{\Theta_{\mathcal{G}}} P(D | \theta_{\mathcal{G}}, \mathcal{G}) P(\theta_{\mathcal{G}} | \mathcal{G}) d\theta_{\mathcal{G}}$$

- There is a closed form solution on page 801 of the book
- If we use a Dirichlet priors for all parameters in the network, as  $M \rightarrow \infty$

$$\log P(D | \mathcal{G}) = l(\hat{\theta}_{\mathcal{G}} : D) - \frac{\log M}{2} \text{Dim}[\mathcal{G}] + O(1)$$

- $l(\theta:D)$  is the loglikelihood and  $\text{Dim}[\mathcal{G}]$  is the model dimension, i.e., the number of independent parameters in  $\mathcal{G}$



# BIC SCORE

$$\begin{aligned} \text{score}_{\text{BIC}}(\mathcal{G} : D) &= l(\hat{\theta}_{\mathcal{G}} : D) - \frac{\log M}{2} \text{Dim}[\mathcal{G}] \\ &= M \sum_{i=1}^n \mathbf{I}_{\hat{P}}(X_i; Pa_{X_i}^{\mathcal{G}}) - M \sum_{i=1}^n \mathbf{H}_{\hat{P}}(X_i) - \frac{\log M}{2} \text{Dim}[\mathcal{G}] \end{aligned}$$

- A trade-off between the fit to the data and the model complexity
- The mutual information grows linearly in  $M$  whereas the model complexity grows logarithmically
  - The larger the  $M$  the more emphasis on fit to the data

# SCORE EQUIVALENCE

$$X \leftarrow Y$$

$$X \rightarrow Y$$

$$X \rightarrow Y \rightarrow Z$$

$$X \leftarrow Y \leftarrow Z$$

$$X \leftarrow Y \rightarrow Z$$

- A score satisfies score equivalence if for all I-Equivalent structures  $G$  and  $G'$ ,  $\text{score}(G : D) = \text{score}(G' : D)$
- The likelihood score satisfies score equivalence
- The Bayesian score satisfies score equivalence if and only if we use BDe priors

# DECOMPOSABLE SCORE

- A structure score  $\text{score}(\mathcal{G} : D)$  is decomposable if it can be written as

$$\text{score}(\mathcal{G} : D) = \sum_i \text{FamScore}(X_i | \text{Pa}_{X_i}^{\mathcal{G}} : D)$$

where  $\text{FamScore}(X | U : D)$  measures how well a set of variables  $U$  serves as parents of  $X$  in  $D$ .

- The likelihood score is decomposable
  - $\text{FamScore}(X | U : D) = M^*[I(X;U) - H(X)]$
- The Bayesian score is decomposable under certain natural assumptions

# WHY IS DECOMPOSABILITY IMPORTANT?


- When searching for structures, small changes to the structure change the score of the structure, but
- The score of the structure can be updated efficiently by recomputing the scores for only the affected families
- For example
  - Adding or deleting an edge changes the score for only one variable
  - Reversing an edge changes the score for only two variables

# STRUCTURE SEARCH

- So far, we have discussed computing a score for a given structure
- How can we find the highest scoring structure?
- Three scenarios
  1. A tree-structured network
  2. Known variable order
  3. General graphs

# TREE-STRUCTURED NETWORKS

- Each variable has at most one parent
- Can be learned efficiently – in polynomial time
- They capture the most important dependencies and they serve as a good baseline
- Define  $\Delta(\mathcal{G}) = \text{score}(\mathcal{G} : D) - \text{score}(\mathcal{G}_{\emptyset} : D)$
- We know  $\text{score}(\mathcal{G}_{\emptyset} : D) = \sum_i \text{FamScore}(X_i : D)$
- $\Delta(\mathcal{G}) = \sum_j \text{FamScore}(X_j \mid \text{Pa}(X_j) : D) - \text{FamScore}(X_j : D)$
- Algorithm
  - Define a fully connected undirected graph where the edge weight between  $X_i$  and  $X_j$  is  $\text{FamScore}(X_i \mid X_j : D) - \text{FamScore}(X_i : D)$
  - Find the maximum spanning forest
- Complexity:  $O(n^2M + n^2 \log n)$



# KNOWN VARIABLE ORDER

- A variable order is given
  - Assuming a variable order is problematic, but in some cases, the order is natural, such as a temporal one
- A variable's parents can be only the variables that precede it

$$Pa_{X_i}^G = \arg \max_{\mathbf{U}_i \subseteq \{X_j : X_j \prec X_i\}} \text{FamScore}(X_i \mid \mathbf{U}_i : D)$$

- To find the parent of the  $n^{\text{th}}$  node, we need to consider  $2^{n-1}$  potential sets
- We can bound the indegree to be  $d$ ; then the algorithm is exponential in  $d$  but polynomial in  $n$

# GENERAL GRAPHS

- Finding the best scoring  $G$  is NP-hard
- Even if we limit the indegree to be  $d$ , finding the best scoring  $G$  is NP-hard when  $d \geq 2$
- Heuristic algorithms that might find the best one but not guaranteed
- This is a combinatorial optimization problem
  - Search the search space to find a good scoring  $G$
  - Use local search
- Need three components
  1. The search space
  2. Scoring function
  3. A search procedure



# SEARCH SPACE

- The search space is a graph
  - Whose nodes are potential solutions,
  - Connected by search operators that allow us to move from one candidate solution to another
- We choose the following search operators
  - Edge addition
  - Edge deletion
  - Edge reversal
- Why can't we just start from  $G_\emptyset$  and use only edge addition operators? Why do we need edge deletion and reversal?
- When is edge reversal useful? Can't we just do it in two steps by deleting an edge and adding back it in the reverse direction?

# SEARCH PROCEDURE

- Start with an initial graph
  - It can be the empty graph, a random graph, the best tree, or a network that is hand constructed
- Apply the search operators to generate new candidate structures, score them, and move to the one that has the highest score
- The runtime is roughly  $O(K * n^2 * (M + n * d))$ 
  - $K$  is the number of iterations,
  - $n$  is the number of variables,
  - $M$  is the number of data instances,
  - $d$  is the max indegree

# LOCAL MAXIMA

- What happens if none of the operators find a better scoring structure than the current one?
  - We hit either a *local maxima* or a *plateau*
- Unfortunately, plateaus are pretty common
  - When the scoring function satisfies score equivalence, then all the I-equivalent structures have equal score
  - The edge reversal operator will generate many I-equivalent structures


# HOW TO DEAL WITH LOCAL MAXIMA AND PLATEAUS?

- *Basin flooding*
  - Keep track of visited states
- *Tabu search*
  - Keep a list of recently performed operations and do not reverse their effect in the next  $L$  steps
- *Data perturbation*
  - Make small changes to data  $D$  through replicas and removals.
  - The new data will cause similar networks to differ, but the big trends will be preserved

## HOW SCORE DECOMPOSITION COMES IN HANDY

- Adding and deleting an edge changes the family of one variable and reversal changes the family of two variables
- We keep track of how much the score would change for each candidate operation
- Let's say we are considering to add the edge  $X \rightarrow Y$ .
  - We compute the score change, but we decide to do a better operation which does not change  $Y$ 's family.
  - In the next step, we do not need to compute the score change for  $X \rightarrow Y$ ; we can use the previous computation.

# HOW GOOD IS LOCAL SEARCH?



# SUMMARY

- Not all the dependencies seen in the data are real
- Determining the edge directionality is sometimes impossible
- We have seen two structure learning approaches
  - Constraint-based
    - Sensitive to independence failures, which are common in real-world datasets
  - Score-based
    - Maximum likelihood score
    - Bayesian Information Criterion (BIC) score
- Tree-structured networks can be learned in polynomial time
- General structure learning is NP-hard but there are many local search algorithms that perform fairly well in practice

# MARKOV NETWORKS



# STRUCTURE LEARNING

- Two approaches
  - Constraint-based
  - Score-based

# CONSTRAINT-BASED

- Assume  $\mathcal{H}^*$  is a P-Map for  $P^*$
- We already know the algorithm
  - See Markov Networks slide deck
  - Use chi2, mutual information for independence test

# SCORE-BASED

- To compute the score for a candidate structure, we need to estimate the parameters
  - Unlike Bayesian networks, parameter estimation cannot be done independently for each variable
  - Unlike Bayesian networks, parameter estimation requires running inference
- Unlike Bayesian networks, the score for a structure does not decompose over variables
  - There is no clean and efficient way of reusing computations

# IN PRACTICE

- Log-linear models are learned
- Search is performed over candidate features (over candidate cliques)
- $L_1$ -regularization is used so that the weights of the useless features become zero
  - If all the weights for all features for a candidate clique become zero, that clique is gone
  - If there is a single feature with a non-zero weight, the clique is retained
  - To achieve sparse solutions, *block  $L_1$ -regularization* is used where a block of features for a candidate clique are penalized together